

Learning to Predict Based on Causality Using an Adversarial Objective

Ryen Krusinga¹, David Jacobs²

¹University of Maryland

²University of Maryland

{krusinga, djacobs}@umiacs.umd.com

Abstract

In a given environment, some variables are causal, and some are merely correlated with the prediction targets. Standard machine learning models make no distinction between the two, making the models less robust to domain shift, in which causal factors remain invariant but correlations change. Given background knowledge about the causal structure of the environment, we demonstrate a simple adversarial method to train a predictor that ignores non-causal information. We show results on two artificial environments, one simple and low-dimensional, the other high-dimensional.

This is a test citation [Oh *et al.*, 2015]. Here is another test sentence. [Oh *et al.*, 2015]. Further test sentences.

References

[Oh *et al.*, 2015] Junhyuk Oh, Xiaoxiao Guo, Honglak Lee, Richard Lewis, and Satinder Singh. Action-Conditional Video Prediction using Deep Networks in Atari Games. pages 1–9, 2015.